

# Part A

## Exercise 1

In kern/pmap.c, mem\_init()

```
////////////////////////////////////
// Make 'envs' point to an array of size 'NENV' of 'struct Env'.
// LAB 3: Your code here.
envs = (struct Env *)boot_alloc(NENV * sizeof(struct Env));
memset(envs, 0, NENV * sizeof(struct Env));

////////////////////////////////////
// Map the 'envs' array read-only by the user at linear address UENVS
// (ie. perm = PTE_U | PTE_P).
// Permissions:
//   - the new image at UENVS -- kernel R, user R
//   - envs itself -- kernel RW, user NONE
// LAB 3: Your code here.
boot_map_region(kern_pgdir, UENVS, PTSIZE, PADDR(envs), PTE_U);
```

## Exercise 2

In kern/env.c

```
// Mark all environments in 'envs' as free, set their env_ids to 0,
// and insert them into the env_free_list.
// Make sure the environments are in the free list in the same order
// they are in the envs array (i.e., so that the first call to
// env_alloc() returns envs[0]).
//
void
env_init(void)
{
    // Set up envs array
    // LAB 3: Your code here.
    for (int i = NENV - 1; i >= 0; i--)
    {
        envs[i].env_id = 0;
        envs[i].env_status = ENV_FREE;
        envs[i].env_link = env_free_list;
        env_free_list = &envs[i];
    }

    // Per-CPU part of the initialization
    env_init_percpu();
}
```

```

static int
env_setup_vm(struct Env *e)
{
    int i;
    struct PageInfo *p = NULL;

    // Allocate a page for the page directory
    if (!(p = page_alloc(ALLOC_ZERO)))
        return -E_NO_MEM;

    // Now, set e->env_pgdir and initialize the page directory.
    //
    // Hint:
    //   - The VA space of all envs is identical above UTOP
    //     (except at UVPT, which we've set below).
    //     See inc/memlayout.h for permissions and layout.
    //     Can you use kern_pgdir as a template? Hint: Yes.
    //     (Make sure you got the permissions right in Lab 2.)
    //   - The initial VA below UTOP is empty.
    //   - You do not need to make any more calls to page_alloc.
    //   - Note: In general, pp_ref is not maintained for
    //     physical pages mapped only above UTOP, but env_pgdir
    //     is an exception -- you need to increment env_pgdir's
    //     pp_ref for env_free to work correctly.
    //   - The functions in kern/pmap.h are handy.

    // LAB 3: Your code here.
    p->pp_ref++;
    e->env_pgdir = page2kva(p);
    for (i = PDX(UTOP); i < NPENTRIES; i++)
    {
        e->env_pgdir[i] = kern_pgdir[i];
    }

    // UVPT maps the env's own page table read-only.
    // Permissions: kernel R, user R
    e->env_pgdir[PDX(UVPT)] = PADDR(e->env_pgdir) | PTE_P | PTE_U;

    return 0;
}

```

```

//
// Allocate len bytes of physical memory for environment env,
// and map it at virtual address va in the environment's address space.
// Does not zero or otherwise initialize the mapped pages in any way.
// Pages should be writable by user and kernel.
// Panic if any allocation attempt fails.
//
void
region_alloc(struct Env *e, void *va, size_t len)
{
    // LAB 3: Your code here.
    // (But only if you need it for load_icode.)
    //
    // Hint: It is easier to use region_alloc if the caller can pass
    // 'va' and 'len' values that are not page-aligned.
    // You should round va down, and round (va + len) up.
    // (Watch out for corner-cases!)
    uintptr_t begin = (uintptr_t)ROUNDDOWN(va, PGSIZE);
    uintptr_t end = (uintptr_t)ROUNDUP(va + len, PGSIZE);
    for (uintptr_t i = begin; i < end; i += PGSIZE)
    {
        struct PageInfo *pp = page_alloc(ALLOC_ZERO);
        if (!pp)
        {
            panic("In region_alloc: failed to allocate page to user environment");
        }
        if (page_insert(e->env_pgdir, pp, (void *)i, PTE_U | PTE_W) != 0)
        {
            panic("In region_alloc: failed to insert page into user environment");
        }
    }
}

```

```

static void
load_icode(struct Env *e, uint8_t *binary)
{
    // LAB 3: Your code here.
    struct Elf *elf = (struct Elf *)binary;
    if (elf->e_magic != ELF_MAGIC)
    {
        panic("In load_icode: invalid ELF");
    }

    struct Proghdr *ph, *eph;
    ph = (struct Proghdr *) (binary + elf->e_phoff);
    eph = ph + elf->e_phnum;
    for (; ph < eph; ph++)
    {
        if (ph->p_type == ELF_PROG_LOAD)
        {
            region_alloc(e, (void *)ph->p_va, ph->p_memsz);
            memset((void *)ph->p_va, 0, ph->p_memsz);
            memmove((void *)ph->p_va, binary+ph->p_offset, ph->p_filesz);
        }
    }

    // set the entry point
    e->env_tf.tf_eip = elf->e_entry;

    // Now map one page for the program's initial stack
    // at virtual address USTACKTOP - PGSIZE.

    // LAB 3: Your code here.
    region_alloc(e, (void *) (USTACKTOP - PGSIZE), PGSIZE);
    e->env_brk = (uintptr_t)ROUNDDOWN(USTACKTOP - PGSIZE, PGSIZE); // update current program's break
}

```

```

//
// Allocates a new env with env_alloc, loads the named elf
// binary into it with load_icode, and sets its env_type.
// This function is ONLY called during kernel initialization,
// before running the first user-mode environment.
// The new env's parent ID is set to 0.
//
void
env_create(uint8_t *binary, enum EnvType type)
{
    // LAB 3: Your code here.
    struct Env *e;
    if (env_alloc(&e, 0) != 0)
    {
        panic("In env_create: failed to allocate environment");
    }
    e->env_type = type;

    // switch page table when change environment
    lcr3(PADDR(e->env_pgdir));
    load_icode(e, binary);
    lcr3(PADDR(kern_pgdir));
}

```

```

void
env_run(struct Env *e)
{
    // Step 1: If this is a context switch (a new environment is running):
    //          1. Set the current environment (if any) back to
    //             ENV_RUNNABLE if it is ENV_RUNNING (think about
    //             what other states it can be in),
    //          2. Set 'curenv' to the new environment,
    //          3. Set its status to ENV_RUNNING,
    //          4. Update its 'env_runs' counter,
    //          5. Use lcr3() to switch to its address space.
    // Step 2: Use env_pop_tf() to restore the environment's
    //          registers and drop into user mode in the
    //          environment.

    // Hint: This function loads the new environment's state from
    //        e->env_tf. Go back through the code you wrote above
    //        and make sure you have set the relevant parts of
    //        e->env_tf to sensible values.

    // LAB 3: Your code here.
    if (curenv != e)
    {
        if (curenv && curenv->env_status == ENV_RUNNING)
        {
            curenv->env_status = ENV_RUNNABLE;
        }
        curenv = e;
        e->env_status = ENV_RUNNING;
        e->env_runs++;
        lcr3(PADDR(e->env_pgdir));
    }
    env_pop_tf(&e->env_tf);
}

```

## Exercise 4

In kern/trapentry.S

```
.text

/*
 * Lab 3: Your code here for generating entry points for the different traps.
 */
// 9 and 15 are reserved
TRAPHANDLER_NOEC(ENTRY_DIVIDE, T_DIVIDE)      // divide error
TRAPHANDLER_NOEC(ENTRY_DEBUG, T_DEBUG)        // debug exception
TRAPHANDLER_NOEC(ENTRY_NMI, T_NMI)            // non-maskable interrupt
TRAPHANDLER_NOEC(ENTRY_BRKPT, T_BRKPT)        // breakpoint
TRAPHANDLER_NOEC(ENTRY_OFLOW, T_OFLOW)        // overflow
TRAPHANDLER_NOEC(ENTRY_BOUND, T_BOUND)        // bounds check
TRAPHANDLER_NOEC(ENTRY_ILLOP, T_ILLOP)        // illegal opcode
TRAPHANDLER_NOEC(ENTRY_DEVICE, T_DEVICE)      // device not available
TRAPHANDLER(ENTRY_DBLFLT, T_DBLFLT)           // double fault
TRAPHANDLER(ENTRY_TSS, T_TSS)                 // invalid task switch segment
TRAPHANDLER(ENTRY_SEGNP, T_SEGNP)             // segment not present
TRAPHANDLER(ENTRY_STACK, T_STACK)            // stack exception
TRAPHANDLER(ENTRY_GPFLT, T_GPFLT)            // general protection fault
TRAPHANDLER(ENTRY_PGFLT, T_PGFLT)            // page fault
TRAPHANDLER_NOEC(ENTRY_FPERR, T_FPERR)        // floating point error
TRAPHANDLER(ENTRY_ALIGN, T_ALIGN)            // alignment check
TRAPHANDLER_NOEC(ENTRY_MCHK, T_MCHK)          // machine check
TRAPHANDLER_NOEC(ENTRY_SIMDERR, T_SIMDERR)    // SIMD floating point error
TRAPHANDLER_NOEC(ENTRY_SYSCALL, T_SYSCALL)   // system call

/*
 * Lab 3: Your code here for _alltraps
 */
.globl _alltraps;
.type _alltraps,@function;
.align 2;
_alltraps:
    pushl %ds
    pushl %es
    pushal
    movl $GD_KD, %eax
    movw %ax, %ds
    movw %ax, %es
    pushl %esp
    call trap
```

In kern/trap.c

Declare entry point functions, then use SETGATE macro to initialize IDT

```

// entry points, 9 and 15 reserved
extern void ENTRY_DIVIDE();      // 0 divide error
extern void ENTRY_DEBUG();      // 1 debug exception
extern void ENTRY_NMI();        // 2 non-maskable interrupt
extern void ENTRY_BRKPT();      // 3 breakpoint
extern void ENTRY_OFLOW();      // 4 overflow
extern void ENTRY_BOUND();      // 5 bounds check
extern void ENTRY_ILLOP();      // 6 illegal opcode
extern void ENTRY_DEVICE();     // 7 device not available
extern void ENTRY_DBLFLT();     // 8 double fault
extern void ENTRY_TSS();        // 10 invalid task switch segment
extern void ENTRY_SEGNP();      // 11 segment not present
extern void ENTRY_STACK();      // 12 stack exception
extern void ENTRY_GPFLT();      // 13 general protection fault
extern void ENTRY_PGFLT();      // 14 page fault
extern void ENTRY_FPERR();      // 16 floating point error
extern void ENTRY_ALIGN();      // 17 alignment check
extern void ENTRY_MCHK();       // 18 machine check
extern void ENTRY_SIMDERR();    // 19 SIMD floating point error
extern void ENTRY_SYSCALL();    // 48 system call

void
trap_init(void)
{
    extern struct Segdesc gdt[];

    // LAB 3: Your code here.
    SETGATE(idt[T_DIVIDE ], 1, GD_KT, ENTRY_DIVIDE , 0);
    SETGATE(idt[T_DEBUG  ], 1, GD_KT, ENTRY_DEBUG  , 0);
    SETGATE(idt[T_NMI     ], 0, GD_KT, ENTRY_NMI     , 0);
    SETGATE(idt[T_BRKPT   ], 1, GD_KT, ENTRY_BRKPT   , 3);
    SETGATE(idt[T_OFLOW   ], 1, GD_KT, ENTRY_OFLOW   , 3);
    SETGATE(idt[T_BOUND   ], 1, GD_KT, ENTRY_BOUND   , 3);
    SETGATE(idt[T_ILLOP   ], 1, GD_KT, ENTRY_ILLOP   , 0);
    SETGATE(idt[T_DEVICE  ], 1, GD_KT, ENTRY_DEVICE  , 0);
    SETGATE(idt[T_DBLFLT  ], 1, GD_KT, ENTRY_DBLFLT  , 0);
    SETGATE(idt[T_TSS     ], 1, GD_KT, ENTRY_TSS     , 0);
    SETGATE(idt[T_SEGNP   ], 1, GD_KT, ENTRY_SEGNP   , 0);
    SETGATE(idt[T_STACK   ], 1, GD_KT, ENTRY_STACK   , 0);
    SETGATE(idt[T_GPFLT   ], 1, GD_KT, ENTRY_GPFLT   , 0);
    SETGATE(idt[T_PGFLT   ], 1, GD_KT, ENTRY_PGFLT   , 0);
    SETGATE(idt[T_FPERR   ], 1, GD_KT, ENTRY_FPERR   , 0);
    SETGATE(idt[T_ALIGN   ], 1, GD_KT, ENTRY_ALIGN   , 0);
    SETGATE(idt[T_MCHK     ], 1, GD_KT, ENTRY_MCHK     , 0);
    SETGATE(idt[T_SIMDERR ], 1, GD_KT, ENTRY_SIMDERR , 0);
    SETGATE(idt[T_SYSCALL ], 1, GD_KT, ENTRY_SYSCALL , 3);

    // Per-CPU setup
    trap_init_percpu();
}

```

# Part B

## Exercise 5&6

In kern/trap.c

```
static void
trap_dispatch(struct Trapframe *tf)
{
    // Handle processor exceptions.
    // LAB 3: Your code here.
    switch (tf->tf_trapno)
    {
        case T_PGFLT:
            page_fault_handler(tf);
            return;
        case T_BRKPT:
            monitor(tf);
            return;
        case T_SYSCALL:
```

## Exercise 7

The entry of syscall in IDT can be seen in the answer of exercise 4.

In kern/trap.c, trap\_dispatch()

```
case T_SYSCALL:
    tf->tf_regs.reg_eax = syscall(tf->tf_regs.reg_eax, tf->tf_regs.reg_edx, tf->tf_regs.reg_ecx, tf->tf_regs.reg_ebx, tf->tf_regs.reg_edi, tf->tf_regs.reg_esi);
    return;
```

In kern/syscall.c

```
// Dispatches to the correct kernel function, passing the arguments.
int32_t
syscall(uint32_t syscallno, uint32_t a1, uint32_t a2, uint32_t a3, uint32_t a4, uint32_t a5)
{
    // Call the function corresponding to the 'syscallno' parameter.
    // Return any appropriate return value.
    // LAB 3: Your code here.

    switch (syscallno)
    {
        case SYS_cputs:
            sys_cputs((char *)a1, (size_t)a2);
            return 0;
        case SYS_cgetc:
            return sys_cgetc();
        case SYS_getenvid:
            return sys_getenvid();
        case SYS_env_destroy:
            return sys_env_destroy((envid_t)a1);
        case SYS_map_kernel_page:
            return sys_map_kernel_page((void *)a1, (void *)a2);
        case SYS_sbrk:
            return sys_sbrk(a1);
        case NSYSCALLS:
        default:
            return -E_INVAL;
    }
}
```

## Exercise 8

In kern/trapentry.S

Add sysenter\_handler

```
.global sysenter_handler
.type sysenter_handler, @function
.align 2
sysenter_handler:
    pushl %esi
    pushl %edi
    pushl %ebx
    pushl %ecx
    pushl %edx
    pushl %eax
    call syscall
    movl %esi, %edx
    movl %ebp, %ecx
    sysexit
```

In kern/trap.c

Add declaration of sysenter handler

```
extern void sysenter_handler();
```

Trap\_init()

```
wrmsr(0x174, GD_KT, 0);           // SYSENTER_CS_MSR
wrmsr(0x175, KSTACKTOP, 0);       // SYSENTER_ESP_MSR
wrmsr(0x176, sysenter_handler, 0); // SYSENTER_EIP_MSR
```

In inc/x86.h

Add implementation of wrmsr

```
#define wrmsr(msr, val1, val2) \
    __asm__ __volatile__ ("wrmsr" \
        : /* no outputs */ \
        : "c" (msr), "a" (val1), "d" (val2))
```

In lib/syscall.c

```
asm volatile(
    // not quite understand, aided by others
    // Store return %esp to %ebp, store return pc to %esi
    "pushl %%esp\n\t"
    "popl %%ebp\n\t"
    "leal after_sysenter_label%, %%esi\n\t" // Use "%=" to generate a unique label number.
    "sysenter\n\t"
    "after_sysenter_label%=: \n\t"
    : "=a" (ret)
    : "a" (num),
      "d" (a1),
      "c" (a2),
      "b" (a3),
      "D" (a4)
    : "cc", "memory");
```



## Exercise 9

In lib/libmain.c, libmain()

```
// set thisenv to point at our Env structure in envs[].
// LAB 3: Your code here.
thisenv = &envs[ENVX(sys_getenvid())];
```

## Exercise 10

In inc/env.h, struct Env

Add a new member to record current program's break

```
uintptr_t env_brk;           // current program's break
```

In kern/syscall.c

```
static int
sys_sbrk(uint32_t inc)
{
    // LAB3: your code here.
    // use region_alloc to allocate memory for environment
    region_alloc(curenv, (void *) (curenv->env_brk - inc), inc);
    curenv->env_brk = (uintptr_t) ROUNDDOWN(curenv->env_brk - inc, PGSIZE);
    return curenv->env_brk;    // return current break rather than previous one
}
```

## Exercise 11

In kern/trap.c, page\_fault\_handler()

Use lowest 2 bits in tf\_cs to check page fault

```
// Handle kernel-mode page faults.

// LAB 3: Your code here.
if (!(tf->tf_cs & 0x3))
{
    panic("In page_fault_handler: kernel page fault");
}
```

In kern/pmap.c

```

int
user_mem_check(struct Env *env, const void *va, size_t len, int perm)
{
    // LAB 3: Your code here.
    // first align start and end address so that every relevant pages can be involved
    uintptr_t begin = (uintptr_t)ROUNDDOWN(va, PGSIZE);
    uintptr_t end = (uintptr_t)ROUNDUP(va+len, PGSIZE);
    perm |= PTE_U;           // guarantee perm includes PTE_U
    for (uintptr_t i = begin; i < end; i += PGSIZE)
    {
        pte_t *pte;
        struct PageInfo *pp = page_lookup(env->env_pgdir, (void *)i, &pte);
        // every page needs to exist, be valid, and meet the requirements
        if (!pp || (*pte & PTE_P) == 0 || ((*pte & perm) == 0 && i >= ULIM))
        {
            user_mem_check_addr = (i == begin) ? (uintptr_t)va : i;
            return -E_FAULT;
        }
    }
    return 0;
}

```

In kern/syscall.c

```

// Print a string to the system console.
// The string is exactly 'len' characters long.
// Destroys the environment on memory errors.
static void
sys_cputs(const char *s, size_t len)
{
    // Check that the user has permission to read memory [s, s+len).
    // Destroy the environment if not.

    // LAB 3: Your code here.
    user_mem_assert(curenv, (void *)s, len, PTE_U);

    // Print the string supplied by the user.
    cprintf("%.s", len, s);
}

```

In kern/kdebug.c, debuginfo\_eip()

```

// Make sure this memory is valid.
// Return -1 if it is not. Hint: Call user_mem_check.
// LAB 3: Your code here.
if (user_mem_check(curenv, usd, sizeof(struct UserStabData), PTE_U))
{
    return -1;
}

stabs = usd->stabs;
stab_end = usd->stab_end;
stabstr = usd->stabstr;
stabstr_end = usd->stabstr_end;

// Make sure the STABS and string table memory is valid.
// LAB 3: Your code here.
if (user_mem_check(curenv, stabs, stab_end - stabs, PTE_U) || user_mem_check(curenv, stabstr, stabstr_end - stabstr, PTE_U))
{
    return -1;
}

```

## Exercise 13

In user/evilhello2.c

```

// Invoke a given function pointer with ring0 privilege, then return to ring3
void ring0_call(void (*fun_ptr)(void)) {
    // Here's some hints on how to achieve this.
    // 1. Store the GDT descriptor to memory (sgdt instruction)
    // 2. Map GDT in user space (sys_map_kernel_page)
    // 3. Setup a CALLGATE in GDT (SETCALLGATE macro)
    // 4. Enter ring0 (lcall instruction)
    // 5. Call the function pointer
    // 6. Recover GDT entry modified in step 3 (if any)
    // 7. Leave ring0 (lret instruction)

    // Hint : use a wrapper function to call fun_ptr. Feel free
    //         to add any functions or global variables in this
    //         file if necessary.

    // Lab3 : Your Code Here
    // step 1, store gdt in e_gdt
    struct Pseudodesc r_gdt;
    sgdt(&r_gdt);

    // step 2, map gdt to vaddr
    int t = sys_map_kernel_page((void* )r_gdt.pd_base, (void* )vaddr);
    if (t < 0)
    {
        cprintf("ring0_call: sys_map_kernel_page failed, %e\n", t);
    }

    // step 3
    // set the address of the callgate
    uint32_t base = (uint32_t)(PGNUM(vaddr) << PTXSHIFT);
    uint32_t index = GD_UD >> 3;
    uint32_t offset = PGOFF(r_gdt.pd_base);

    gdt = (struct Segdesc*)(base+offset);
    entry = gdt + index;
    old= *entry;

    // set up callgate
    // put call_fun_ptr (wrapper function) into entry, set privilege level to 3
    SETCALLGATE(*(struct Gatedesc*)entry), GD_KT, call_fun_ptr, 3);

    // step 4 and step 5
    asm volatile("lcall $0x20, $0");
}

```

```

void call_fun_ptr()
{
    evil();

    // step 6 below
    *entry = old;

    // step 7 below
    asm volatile("leave");
    asm volatile("lret");
}

```